

Detect and Respond? Cool Story — Or Just Don't Let the Bad Stuff Start.

Real-world Kubernetes enforcement with Kyverno and KubeArmor

Matt Brown

Github: [sf-matt](#)



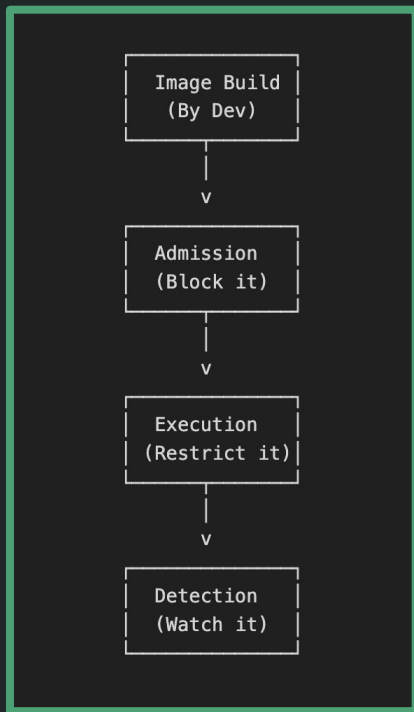
Kubernetes in Plain English

- **Container** - Your app, just the thing you built
- **Pod** - The basic unit in Kubernetes running your containers
- **Workload** - The manager of pods, often a deployment
- **Node** - A machine (virtual or physical) that runs your workloads
- **Cluster** - A group of nodes working together as one system

Kubernetes Insecurity in Plain English

- **Container** - Runs as root, includes package managers, not hardened
- **Pod** - Inherits container defaults, no restrictions applied
- **Workload** - Runs pods as defined, no questions asked
- **Node** - Let's not go there
- **Cluster** - Let's really not go there

Stages Where We Can Secure a Workload



- **Image Build:** Dev and App Sec owned, not Kubernetes managed
- **Admission:** Block bad config before it runs
- **Execution:** Limit what it can do
- **Detection:** Watch what it actually does

Dockerfile — Looks Harmless, Runs as Root

```
FROM python:3.9-slim-buster

WORKDIR /app

COPY requirements.txt requirements.txt
RUN pip3 install -r requirements.txt

COPY . .

CMD ["python3", "-m", "flask", "run", "--host=0.0.0.0"]
```

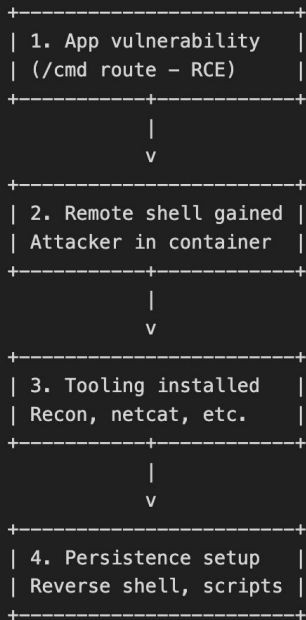
- Based on official image: **python:3.9-slim-buster**
- No **USER** directive → container runs as **UID 0 (root)**
- Works fine — but runs with full privileges by default

Deployment Spec — Kubernetes Doesn't Stop It

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: flask-root
spec:
  replicas: 1
  selector:
    matchLabels:
      app: flask-app
      variant: root
  template:
    metadata:
      labels:
        app: flask-app
        variant: root
    spec:
      containers:
        - name: flask
          image: sf-matt/flask-app:latest
          ports:
            - containerPort: 5000
```

- No **securityContext** → container inherits root from the base image
- No enforcement → container can access files, run commands, and act freely

Shells Happen: From Pod to Compromise



- Default base image (**python:3.9-slim**) runs as root
- Flask app exposes a command injection route (**/cmd**)
- Kubernetes runs it without question
- Attacker gets a remote shell inside the container

curl → connect → compromise

```
matt.brown@matt ~ % curl "http://192.168.64.7:30080/cmd" \  
  --get \  
  --data-urlencode 'input=python3 -c "import socket,os,pty;"
```



Run Payload via RCE

```
matt.brown@matt ~ % nc -l 4444  
# whoami  
whoami  
root  
# apt update -qq && apt install -y -qq nmap  
apt update -qq && apt install -y -qq nmap  
24 packages can be upgraded. Run 'apt list --upgradable' to see them.  
The following additional packages will be installed:  
  libblas3 libgfortran5 liblinear3 liblua5.3-0 libpcap0.8 libssh2-1  
  nmap-common
```



Root access

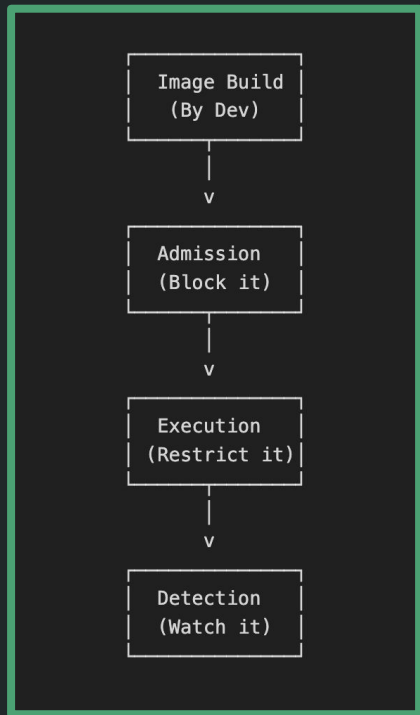


Install recon tooling

Welcome to the **root** Jungle

- **Install Anything:** `apt update && apt install nmap -y`
- **Read Sensitive Files:** `/etc/shadow`
- **Bypass Restrictions:** Override file permissions, explore the filesystem
- **Steal Credentials:** Read Kubernetes service account tokens and secrets from mounted paths
- **Tamper with the Host** (with the right mounts): Use `nsenter` to escape the container

Why Not Just Block It?



- **Image Build:** Dev and App Sec owned, not Kubernetes managed
- **Admission:** Block bad config before it runs
- **Execution:** Limit what it can do
- **Detection:** Watch what it actually does

Kyverno: Kubernetes-Native Admission Control



- Open Source with over 6.4k GitHub stars: <https://github.com/kyverno/kyverno>
- Easy YAML syntax
- Blocks or audits insecure workloads at admission with room for exceptions
- Installs in minutes

What Is an Admission Controller?

- **Admission controllers** run *after* authentication, but *before* persistence
- They mutate or validate the resource
- Validation allows block, audit, pass
- Natively available

This **ClusterPolicy** Blocks Our Flask App!

```
apiVersion: kyverno.io/v1
kind: ClusterPolicy
metadata:
  name: disallow-run-as-root-deployments
spec:
  validationFailureAction: enforce
  rules:
    - name: require-run-as-non-root
      match:
        resources:
          kinds:
            - Deployment
      validate:
        message: "Containers must not run as root (runAsNonRoot: true)."
        pattern:
          spec:
            template:
              spec:
                containers:
                  - securityContext:
                      runAsNonRoot: true
```

```
service/flask-service created
Error from server: error when creating "flask-vuln-demo/flask-deployment.yaml": admission webhook "disallow-run-as-root.kyverno.io" denied the request: resource Deployment/default/flask-app was blocked due to the following policies

disallow-run-as-root:
  autogen-require-run-as-non-root: 'validation error: Containers must not run as root (UID 0). Set runAsNonRoot: true and use a non-root UID, rule autogen-require-run-as-non-root failed at path /spec/template/spec/containers/0/securityContext/'
```

ClusterPolicy Applies Across the Cluster

Set **enforce** to actively block

Apply to any **Deployment**

Clear error **message**

Enforces **runAsNonRoot: true**

Blocked!

Here's What a Compliant **Deployment** Looks Like

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: flask-app
5  spec:
6    replicas: 1
7    selector:
8      matchLabels:
9        app: flask-app
10   template:
11     metadata:
12       labels:
13         app: flask-app
14     spec:
15       containers:
16         - name: flask
17           image: sfmatt/flask-vuln-demo-nonroot:latest
18           ports:
19             - containerPort: 5000
20           securityContext:
21             runAsNonRoot: true
22             runAsUser: 101
```

- Passes Kyverno policy validation with no issues
- Guides developers toward secure-by-default configurations
- Easy to validate in CI pipelines or with **kyverno apply --audit**



Sets **runAsNonRoot** and **runAsUser**

Congrats, your pods don't run as root.

curl → connect → compromise (again?)

```
matt.brown@matt ~ % curl "http://192.168.64.7:30080/cmd" \
--get \
--data-urlencode 'input=python3 -c "import socket,os,pty;
```



Run Payload via RCE

```
matt.brown@matt ~ % nc -l 4444
$ whoami
whoami
appuser
$ apt update -qq && apt install -y -qq nmap
apt update -qq && apt install -y -qq nmap
E: List directory /var/lib/apt/lists/partial is missing. - Acquire (13: P
$ echo "test" > /tmp/hello.txt
echo "test" > /tmp/hello.txt
$ cat /tmp/hello.txt
cat /tmp/hello.txt
test
$ python3 -m http.server 8888
python3 -m http.server 8888
_Serving HTTP on 0.0.0.0 port 8888 (http://0.0.0.0:8888/) ...
```



No root access



No recon tooling




Write to **/tmp**



Python server on high
port

Why Not Just Limit It?



- **Image Build:** Dev and App Sec owned, not Kubernetes managed
- **Admission:** Block bad config before it runs
- **Execution:** Limit what it can do
- **Detection:** Watch what it actually does



KubeArmor: Execution Enforcement

- Open Source with 1.7k+ GitHub stars: <https://github.com/kubearmor/KubeArmor>
- Controls execution via Linux Security Modules (e.g. AppArmor and SELinux)
- Easy YAML syntax
- Abstracts away Linux distro differences

Linux Security Modules (LSMs)

- Part of Linux Kernel
- LSMs intercept **syscalls** and decide whether to allow them
- AppArmor and SELinux are the most widely used (and supported by KubeArmor)
- KubeArmor generates and applies LSM profiles via YAML

KubeArmor in Action: Blocking the Next Step

```
1  apiVersion: security.kubearmor.com/v1
2  kind: KubeArmorPolicy
3  metadata:
4    name: block-tmp-write
5  spec:
6    selector:
7      matchLabels:
8        app: flask-app
9    file:
10     matchDirectories:
11       - dir: /tmp/
12         recursive: true
13     action: Block
14     severity: 5
15 ---
16 apiVersion: security.kubearmor.com/v1
17 kind: KubeArmorPolicy
18 metadata:
19   name: block-python-listener
20 spec:
21   selector:
22     matchLabels:
23       app: flask-app
24   process:
25     matchPaths:
26       - path: /usr/local/bin/python3.9
27     action: Block
28     severity: 5
```



KubeArmorPolicy



Applies to everything done in `/tmp`



Set `action` to explicitly `Block`



Applies to `python3.9` process



Set `action` to explicitly `Block`

curl → no connect → no compromise

```
matt.brown@matt ~ % curl "http://192.168.64.7:30080/cmd" \  
--get \  
--data-urlencode 'input=python3 -c "import socket,os,pty:'
```



Run Payload via RCE

```
matt@controlplane: /kubernetes/podexec$ ca  
matt@controlplane:~$ kubectl logs flask-app-6fd9dbd494-zbfvg  
* Debug mode: off  
WARNING: This is a development server. Do not use it in a production  
* Running on all addresses (0.0.0.0)  
* Running on http://127.0.0.1:5000  
* Running on http://10.244.49.81:5000  
Press CTRL+C to quit  
192.168.64.4 - - [16/Jun/2025 19:50:32] "GET / HTTP/1.1" 200 -  
192.168.64.4 - - [16/Jun/2025 19:50:32] "GET /favicon.ico HTTP/1.1"  
Traceback (most recent call last):  
  File "<string>", line 1, in <module>  
TimeoutError: [Errno 110] Connection timed out  
192.168.64.4 - - [16/Jun/2025 19:51:38] "GET /cmd?input=python3+-c+"  
h\\")" HTTP/1.1" 200 -  
sh: 1: python3: Permission denied  
192.168.64.4 - - [17/Jun/2025 06:18:39] "GET /cmd?input=python3+-c+"  
h\\")" HTTP/1.1" 200 -  
192.168.64.4 - - [17/Jun/2025 06:19:41] "GET / HTTP/1.1" 200 -
```

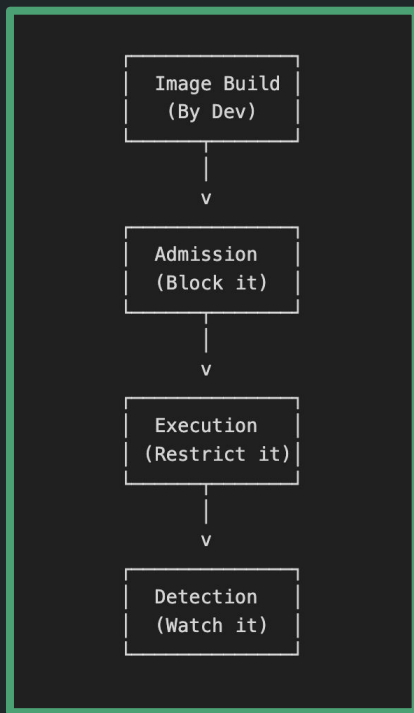


No Reverse Shell



App Accessible

We Can't Block Everything



- **Image Build:** Dev and App Sec owned, not Kubernetes managed
- **Admission:** Block bad config before it runs
- **Execution:** Limit what it can do
- **Detection:** Watch what it actually does

Where Prevention Hits a Wall and eBPF Detection Step In

Prevention Has Limits

- Some pods need to run as root (e.g. monitoring)
- Legit containers may execute perceived malicious behavior (e.g. install packages)
- Scoped policies reduce blast radius but create blind spots

Detection Steps In

- Watch what root does in runtime
- Detect abnormal binary execution, DNS traffic, other syscall activity
- Alert on violations even in "trusted" images

Block Early, Control Behavior, Watch What Matters

1. **Block Early** — Admission control with **Kyverno** prevents insecure workloads from ever starting
2. **Control Behavior** — **KubeArmor** enforces strict LSM policies to stop abuse before it spreads
3. **Watch What Matters** — **eBPF detection agents** catches the unexpected, even in “compliant” pods

Secure Workloads, No Amazon Packages Left To Chance



Hit me up after or online



cloudsecburrito.com



github.com/sf-matt/k8s-enforcement-lab